



# Deep Learning


CNNs






# How do Computers see?







# How do Humans see?





# Mammalian Visual System



# Mammalian Visual System

Cones and rods pick up photons and propagate the signal to the back of our brain

Along the way, the signal gets processed in stages





# Mammalian Visual System

Cones and rods pick up photons and propagate the signal to the back of our brain

Along the way, the signal gets processed in stages

The extracted information gets aggregated and formed into an image





# Mammalian Visual System

Cones and rods pick up photons and propagate the signal to the back of our brain


Along the way, the signal gets processed in stages

The extracted information gets aggregated and formed into an image





# Convolutional Neural Networks




High Performance Machine Learning Group





# Image Representation

01. Is a matrix or grid of intensity values
02. integers  $[0, 255]$  or float points  $[0, 1]$
03. Each element in the matrix is a **pixel**
04. Can have 1 greyscale channel or multiple colour channels: **RGB**






# Image Representation

01. Is a matrix or grid of intensity values
02. integers  $[0, 255]$  or float points  $[0, 1]$
03. Each element in the matrix is a **pixel**
04. Can have 1 greyscale channel or multiple colour channels: **RGB**




# Image Representation

01. Is a matrix or grid of intensity values
02. integers  $[0, 255]$  or float points  $[0,1]$
03. Each element in the matrix is a **pixel**
04. Can have 1 greyscale channel or multiple colour channels: **RGB**



# Correlated structures




Images are **correlated spatially**

Strong correlation between  
neighbouring pixels

Shifting an object will not alter its  
essence

# Correlated structures





Images are **correlated spatially**

Strong correlation between  
neighbouring pixels

Shifting an object will not alter its  
essence



# Correlated structures





Images are **correlated spatially**

Strong correlation between  
neighbouring pixels

Shifting an object will not alter its  
essence

# Correlated structures






Images are **correlated spatially**

Strong correlation between  
neighbouring pixels

Shifting an object will not alter its  
essence

# Correlated structures






Images are **correlated spatially**

Strong correlation between  
neighbouring pixels

Shifting an object will not alter its  
essence



# Correlated structures



Detect the features across an image and aggregate them

Form complete features higher up in the **hierarchy**

Images are **correlated spatially**

Strong correlation between neighbouring pixels

Shifting an object will not alter its essence

# ImageNet Challenge

Computer Vision Benchmark

1.4M Images, 1000 classes

Image Classification

Difficult until 2012



# ImageNet Challenge


Computer Vision Benchmark

1.4M Images, 1000 classes

Image Classification

Difficult until 2012

ImageNet Classification Error (Top 5)









# Convolutions

**Kernel** (filter): small matrix that we use to convolve an image

## Convolution:

An operation that “blends” one function with another.

Operation	Kernel	Image result
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	




# Convolutions

**Kernel** (filter): small matrix that we use to convolve an image

## Convolution:

An operation that “blends” one function with another.

A filter applies a **convolution** operation on an image



Operation	Kernel	Image result
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	



# Convolutions

**Kernel** (filter): small matrix that we use to convolve an image

## Convolution:

An operation that “blends” one function with another.

A filter applies a **convolution** operation on an image





# Convolutions

**Kernel** (filter): small matrix that we use to convolve an image

## Convolution:

An operation that “blends” one function with another.

A filter applies a **convolution** operation on an image



# Convolutions Compute

2	2	6	8	2	0	1	2
4	3	4	5	1	9	6	2
3	9	4	4	7	7	6	9
1	3	4	6	8	2	2	1
8	4	6	2	3	1	8	8
5	8	9	0	1	0	2	2
9	2	6	6	3	6	2	1
9	8	8	2	6	3	4	5


Source layer (image)

Convolutional  
Kernel (filter)

-1	0	1
-2	0	2
-1	0	1



Destination layer


$$\begin{aligned} &(-1 \times 2) + (0 \times 2) + (1 \times 6) + \\ &(-2 \times 4) + (0 \times 3) + (2 \times 4) + \\ &(-1 \times 3) + (0 \times 9) + (1 \times 4) = 5 \end{aligned}$$







# Convolutions Compute



The kernel is **shifted** across the image and produces a point value  
The step size in which it shifts is the **stride**  
The output is always smaller, we use **padding** to preserve dimensions

# Convolutions Compute



The kernel is **shifted** across the image and produces a point value  
The step size in which it shifts is the **stride**  
The output is always smaller, we use **padding** to preserve dimensions

The image shrinks according to  
 $\text{Output\_size} = \text{inputSize} - (\text{KernelSize} - 1)$


# Convolutions

Source layer (image)

2	2	6	8	2	0	1	2
4	3	4	5	1	9	6	3
3	9	4	4	7	7	6	9
1	3	4	6	8	2	2	1
8	4	6	2	3	1	8	8
5	8	9	0	1	0	2	3
9	2	6	6	3	6	2	1
9	8	8	2	6	3	4	5

Feature map (activation map)

	5						



$W_1$	$W_2$	$W_3$
$W_4$	$W_5$	$W_6$
$W_7$	$W_8$	$W_9$

Kernel (filter)


# Convolutions

Source layer (image)

2	2	6	8	2	0	1	2
4	3	4	5	1	9	6	3
3	9	4	4	7	7	6	9
1	3	4	6	8	2	2	1
8	4	6	2	3	1	8	8
5	8	9	0	1	0	2	3
9	2	6	6	3	6	2	1
9	8	8	2	6	3	4	5

Feature map (activation map)

	5						



$W_1$	$W_2$	$W_3$
$W_4$	$W_5$	$W_6$
$W_7$	$W_8$	$W_9$

Kernel (filter)

# Convolutions

Source layer (image)

2	2	6	8	2	0	1	2
4	3	4	5	1	9	6	3
3	9	4	4	7	7	6	9
1	3	4	6	8	2	2	1
8	4	6	2	3	1	8	8
5	8	9	0	1	0	2	3
9	2	6	6	3	6	2	1
9	8	8	2	6	3	4	5

Feature map (activation map)

	5						



$W_1$	$W_2$	$W_3$
$W_4$	$W_5$	$W_6$
$W_7$	$W_8$	$W_9$

Kernel (filter)



How do we know which kernels to use?

Kernels are learnt and initialized randomly during training the CNN learns spatial features

# Convolutions

Source layer (image)

2	2	6	8	2	0	1	2
4	3	4	5	1	9	6	3
3	9	4	4	7	7	6	9
1	3	4	6	8	2	2	1
8	4	6	2	3	1	8	8
5	8	9	0	1	0	2	3
9	2	6	6	3	6	2	1
9	8	8	2	6	3	4	5

Feature map (activation map)

	5						

$w_1$	$w_2$	$w_3$
$w_4$	$w_5$	$w_6$
$w_7$	$w_8$	$w_9$

Kernel (filter)

How do we know which kernels to use?

Kernels are learnt and initialized randomly during training the CNN learns spatial features

Fully connected

$$\sum_{i \in \text{image}}^{W \times H \times C} \mathbf{x}_i \mathbf{w}_i$$

# Convolutions

Source layer (image)

2	2	6	8	2	0	1	2
4	3	4	5	1	9	6	3
3	9	4	4	7	7	6	9
1	3	4	6	8	2	2	1
8	4	6	2	3	1	8	8
5	8	9	0	1	0	2	3
9	2	6	6	3	6	2	1
9	8	8	2	6	3	4	5

Feature map (activation map)

	5						

$w_1$	$w_2$	$w_3$
$w_4$	$w_5$	$w_6$
$w_7$	$w_8$	$w_9$

Kernel (filter)

How do we know which kernels to use?

Kernels are learnt and initialized randomly during training the CNN learns spatial features

Fully connected

$$\sum_{i \in \text{image}} \mathbf{x}_i \mathbf{w}_i \longrightarrow$$

# Convolutions

Source layer (image)

2	2	6	8	2	0	1	2
4	3	4	5	1	9	6	3
3	9	4	4	7	7	6	9
1	3	4	6	8	2	2	1
8	4	6	2	3	1	8	8
5	8	9	0	1	0	2	3
9	2	6	6	3	6	2	1
9	8	8	2	6	3	4	5

Feature map (activation map)

	5						

$W_1$	$W_2$	$W_3$
$W_4$	$W_5$	$W_6$
$W_7$	$W_8$	$W_9$

Kernel (filter)

How do we know which kernels to use?

Kernels are learnt and initialized randomly during training the CNN learns spatial features

Fully connected


$$\sum_{i \in \text{image}}^{W \times H \times C} \mathbf{x}_i \mathbf{w}_i$$

Locally connected shared weights

$$\sum_{i \in 3 \times 3}^{W \times H \times C} \mathbf{x}_i \mathbf{w}_i$$




# Anatomy of an Convolutional Neural Network



CNNs actually look a lot like normal Dense Neural Networks

**Interpretation:** Every node is a kernel

# Anatomy of an Convolutional Neural Network




See the nodes as moving across the image!

We say these are **shared weights**

CNNs actually look a lot like normal Dense Neural Networks

**Interpretation:** Every node is a kernel

# Anatomy of an Convolutional Neural Network



Fully connected

$$\sum_{i \in \text{image}}^{W \times H \times C} \mathbf{x}_i \mathbf{w}_i$$



Locally connected  
shared weights

$$\sum_{i \in 3 \times 3}^{W \times H \times C} \mathbf{x}_i \mathbf{w}_i$$


See the nodes as moving across the image!

We say these are **shared weights**

CNNs actually look a lot like normal Dense Neural Networks

**Interpretation:** Every node is a kernel

# Anatomy of an Convolutional Neural Network



Fully connected

$$\sum_{i \in \text{image}}^{W \times H \times C} \mathbf{x}_i \mathbf{w}_i$$



Locally connected  
shared weights

$$\sum_{i \in 3 \times 3}^{W \times H \times C} \mathbf{x}_i \mathbf{w}_i$$

$w_1$	$w_2$	$w_3$
$w_4$	$w_5$	$w_6$
$w_7$	$w_8$	$w_9$


See the nodes as moving across the image!

We say these are **shared weights**

CNNs actually look a lot like normal Dense Neural Networks



**Interpretation:** Every node is a kernel

# Pooling and dilated convolutions




We could discard information gradually by max pooling  
Keep the strongest signal, works better than average pooling  
Nowadays, we use a bigger stride for dimensionality reduction

# Pooling and dilated convolutions




We could discard information gradually by max pooling  
Keep the strongest signal, works better than average pooling  
Nowadays, we use a bigger stride for dimensionality reduction


Dilated convolutions can be used if you expect your images to have information which are spatially far from each other



# CNN Cases




# AlexNet (2012)







# AlexNet (2012)




Images become smaller and number of filters increases as we go deeper  
8 layers with dropout and ReLU activations trained for 6 days on 2 GPUs

Type: ML Perceptron  
Data Set: MNIST  
Hidden Layers: 3  
Hidden Neurons: 10000  
Synapses: 24864180  
Synapses shown: 12%  
Learning: BP








# Transfer Learning








# Transfer Learning






# Thank You





# VGG16 and ResNet (2014, 2015)



VGG16: construct large and deep models (120M params)

Vanishing gradient problem

# VGG16 and ResNet (2014, 2015)


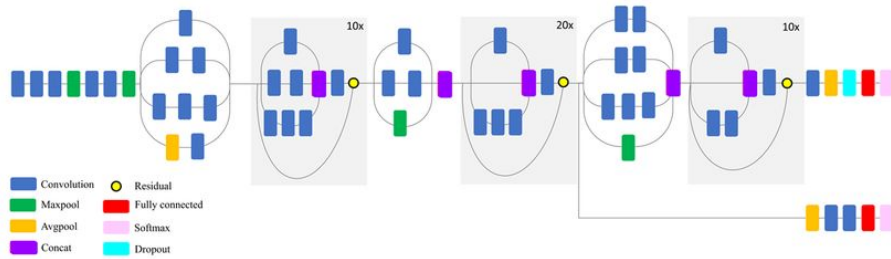


VGG16: construct large and deep models (120M params)


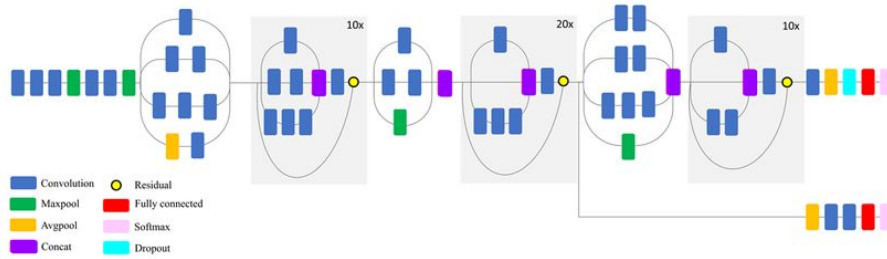
Vanishing gradient problem



# ResNet (2015)




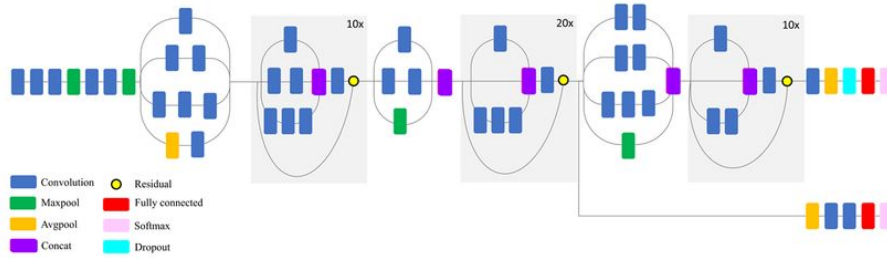
# ResNet (2015)



ResNets: construct large and deep models with skip connections, able to train up to 152 layers (!)

Higher abstraction and less nuisance from vanishing or exploding gradients

# ResNet (2015)





ResNets: construct large and deep models with skip connections, able to train up to 152 layers (!)

Higher abstraction and less nuisance from vanishing or exploding gradients

Early layers learn features that get progressively more abstract, we can **preserve** and **control** the flow of information with **skip connections**

# DenseNet (2016)





ResNets: construct large and deep models with skip connections, able to train up to 152 layers (!)

Higher abstraction and less nuisance from vanishing or exploding gradients

Early layers learn features that get progressively more abstract, we can **preserve** and **control** the flow of information with **skip connections**

# DenseNet (2016)



ResNets: construct large and deep models with skip connections, able to train up to 152 layers (!)

Higher abstraction and less nuisance from vanishing or exploding gradients

Early layers learn features that get progressively more abstract, we can **preserve** and **control** the flow of information with **skip connections**