09/11/2022

Deep Learning

CNNs

Bryan Cardenas
Robert Jan Schlimbach
Caspar van leeuwen

High Performance Machine Learning Group

SURF

How do Computers see?

How do Humans see?

Mammalian Visual System

Cones and rods picks up photons and propagate the signal to the back of our brain

Along the way, the signal gets processed in stages

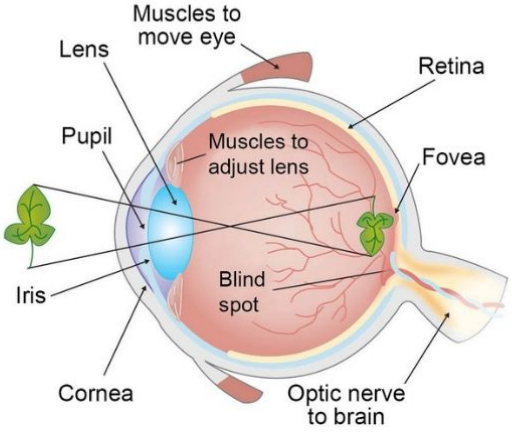

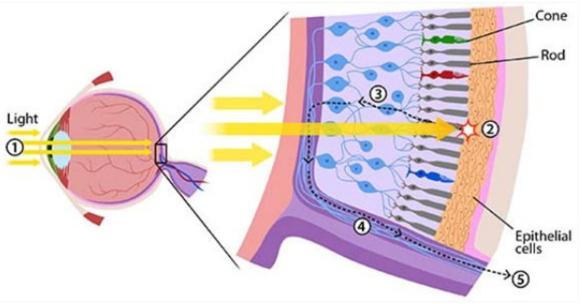

SURF

Mammalian Visual System

Cones and rods picks up photons and propagate the signal to the back of our brain

Along the way, the signal gets processed in stages

The extracted information gets aggregated and formed into an image

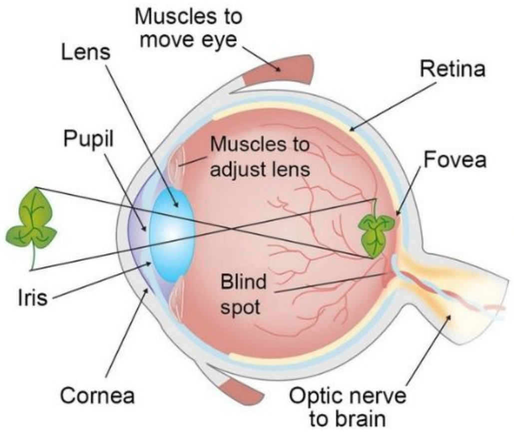

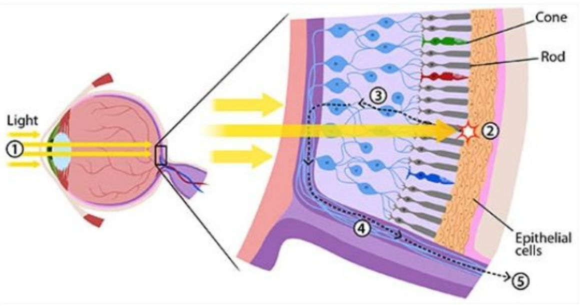

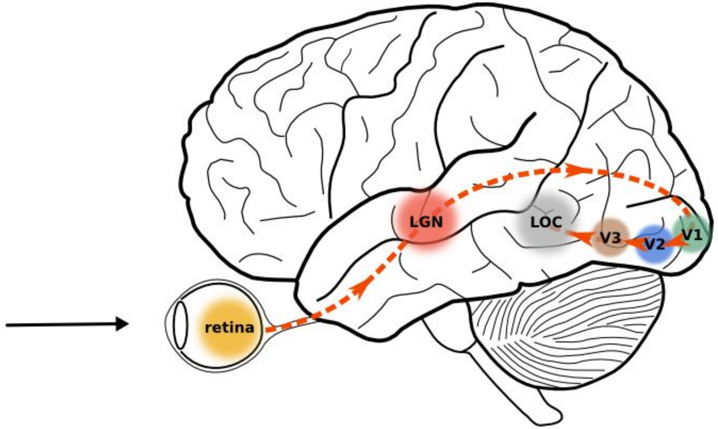

Mammalian Visual System

Cones and rods picks up photons and propagate the signal to the back of our brain

Along the way, the signal gets processed in stages

The extracted information gets aggregated and formed into an image

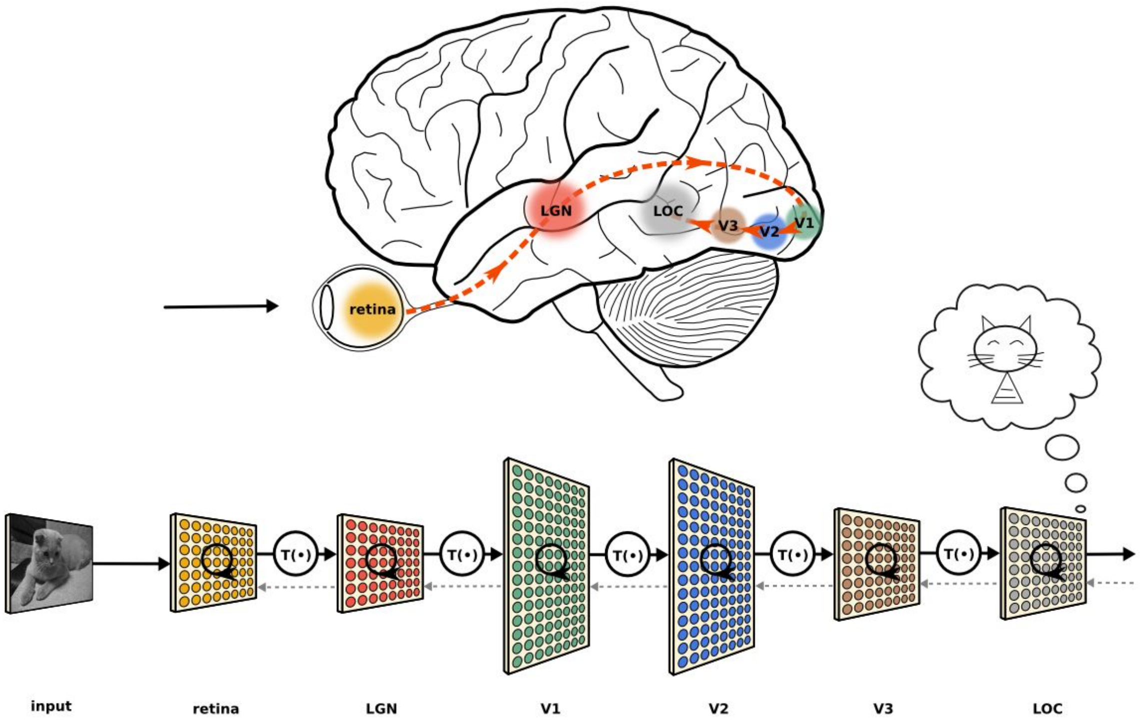

SURF

02. Convolutional Neural Networks

SURF

Image Representation

01. Is a matrix or grid of intensity values

02. integers [0, 255] or float points [0,1]

03. Each element in the matrix is a **pixel**

04. Can have 1 greyscale channel or multiple colour **channels: RGB**

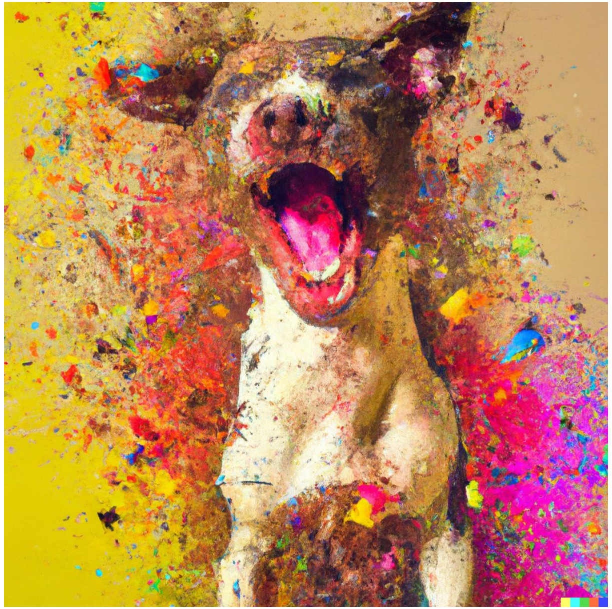

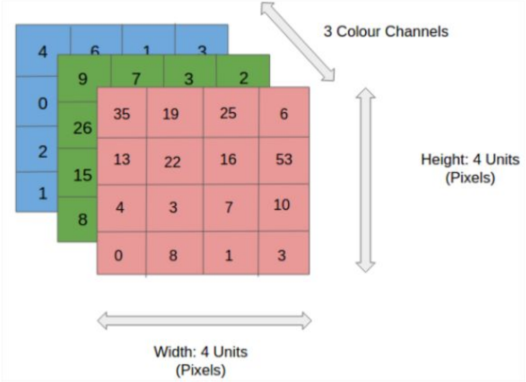

SURF

Image Representation

01. Is a matrix or grid of intensity values

02. integers [0, 255] or float points [0,1]

03. Each element in the matrix is a **pixel**

04. Can have 1 greyscale channel or multiple colour **channels: RGB**

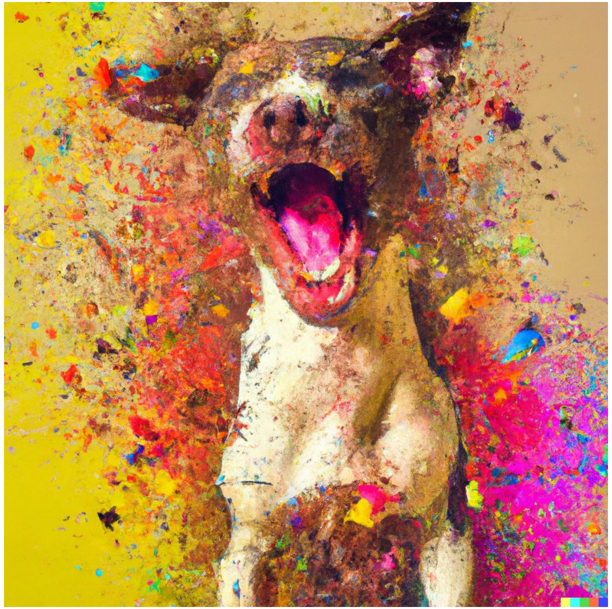

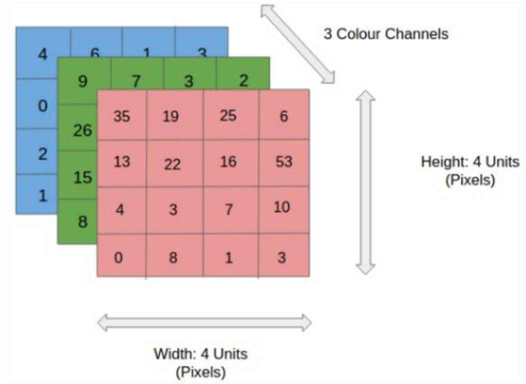

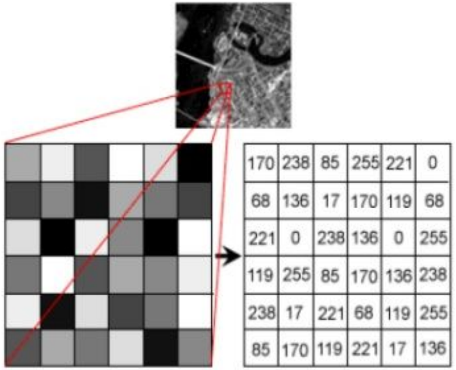

SURF

Correlated structures

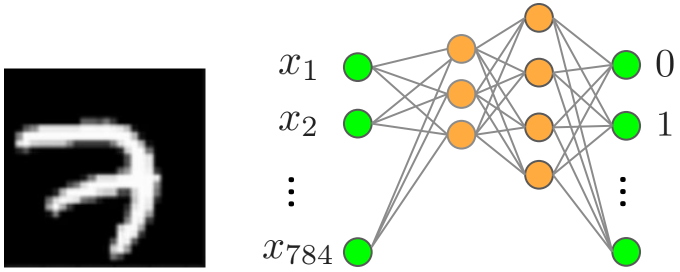

Images are **correlated spatially**

Strong correlation between neighbouring pixels

Shifting an object will not alter its essence

Correlated structures

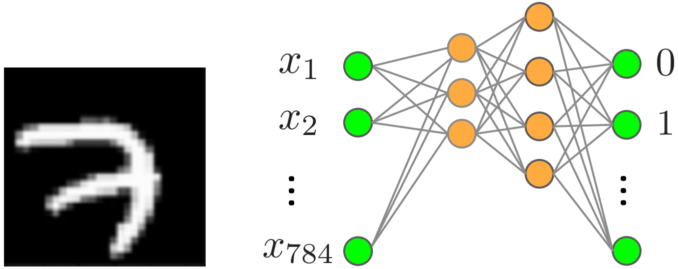

Images are **correlated spatially**

Strong correlation between neighbouring pixels

Shifting an object will not alter its essence

Correlated structures

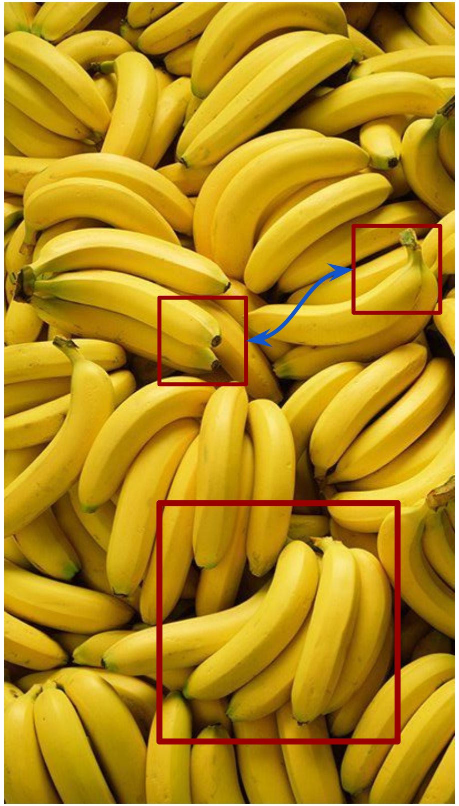

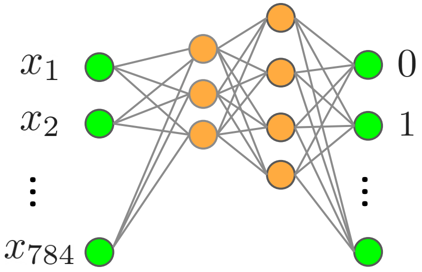

Images are **correlated spatially**

Strong correlation between neighbouring pixels

Shifting an object will not alter its essence

Correlated structures

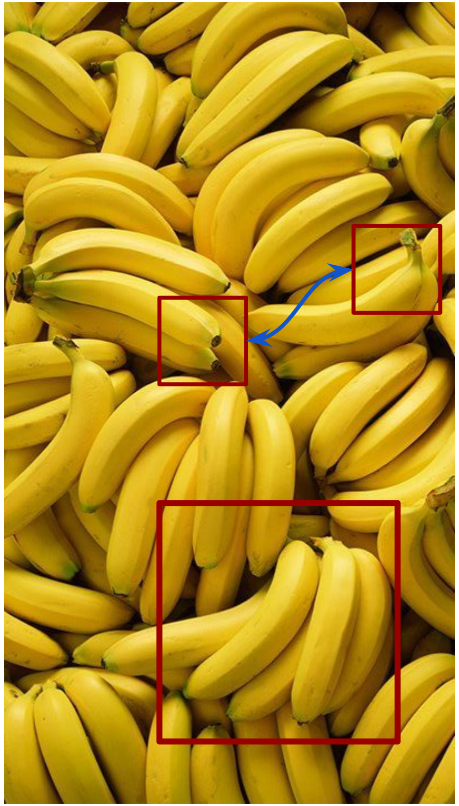

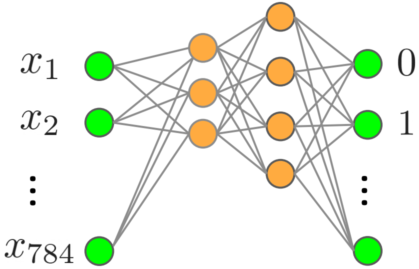

Images are **correlated spatially**

Strong correlation between neighbouring pixels

Shifting an object will not alter its essence

Correlated structures

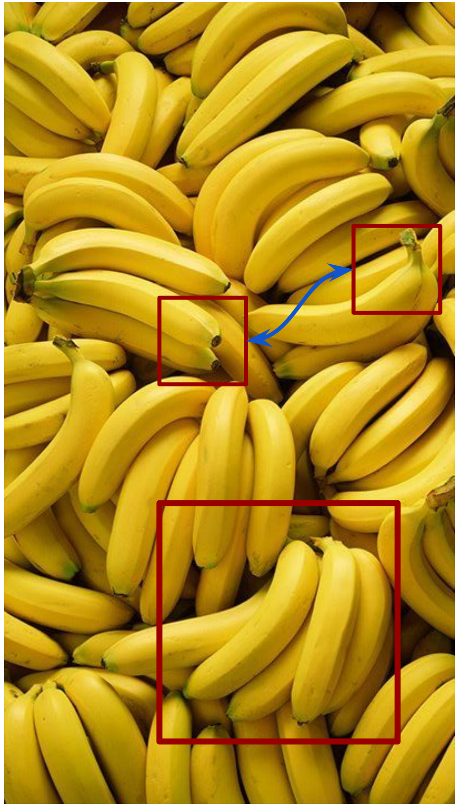

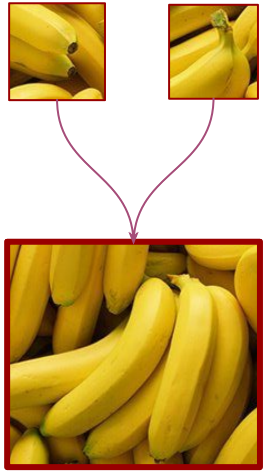

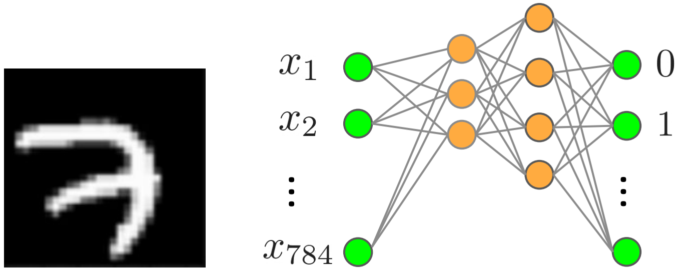

Images are **correlated spatially**

Strong correlation between neighbouring pixels

Shifting an object will not alter its essence

Correlated structures

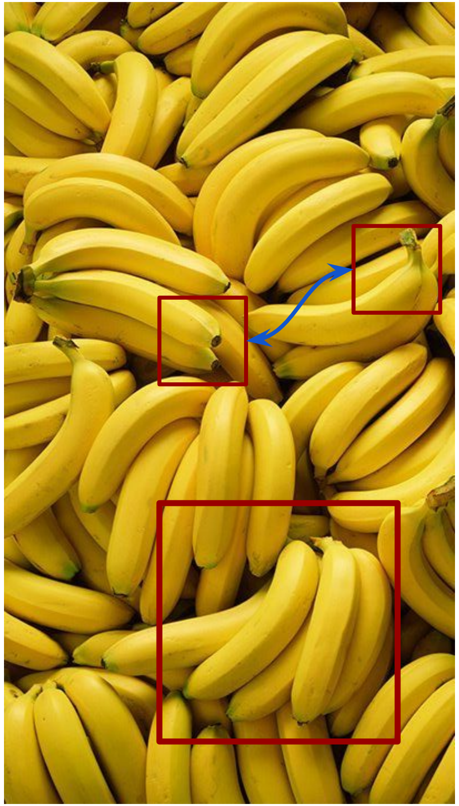

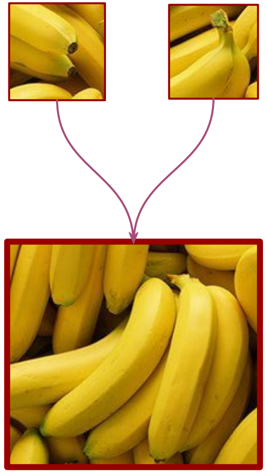

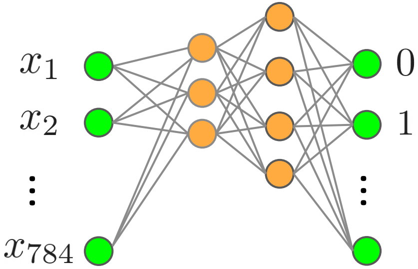

Detect the features across an image and aggregate them

Form complete features higher up in the **hierarchy**

Images are **correlated spatially**

Strong correlation between neighbouring pixels

Shifting an object will not alter its essence

ImageNet Challenge

Computer Vision Benchmark	1.4M Images, 1000 classes
Image Classification	Difficult until 2012

SURF

ImageNet Challenge

Computer Vision Benchmark

1.4M Images, 1000 classes

Image Classification

Difficult until 2012

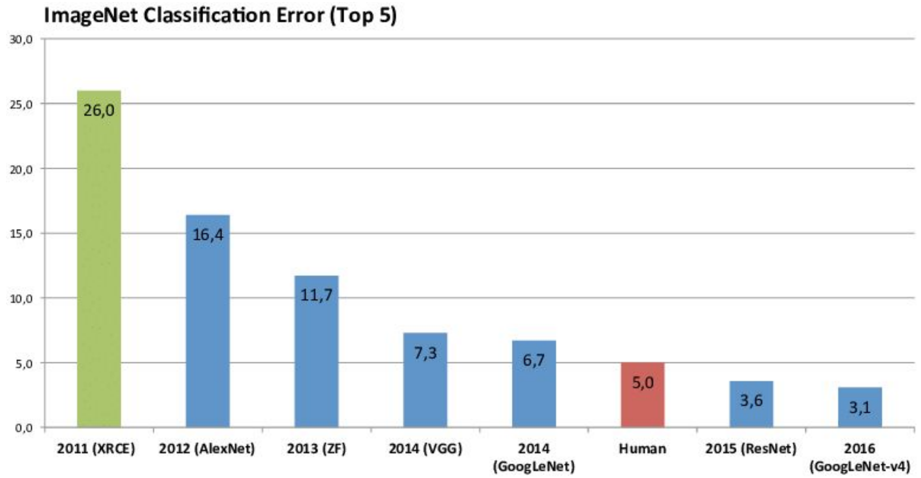

Convolutions

Kernel (filter): small matrix that we use to convolve an image

Convolution:

An operation that “blends” one function with another.

Operation	Kernel	Image result
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9}\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	

Convolutions

Kernel (filter): small matrix that we use to convolve an image

Convolution:

An operation that “blends” one function with another.

A filter applies a **convolution** operation on an image

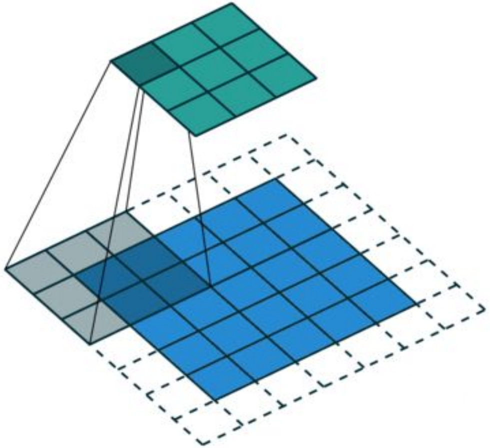

Operation	Kernel	Image result
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9}\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	

Convolutions

Kernel (filter): small matrix that we use to convolve an image

Convolution:

An operation that “blends” one function with another.

A filter applies a **convolution** operation on an image

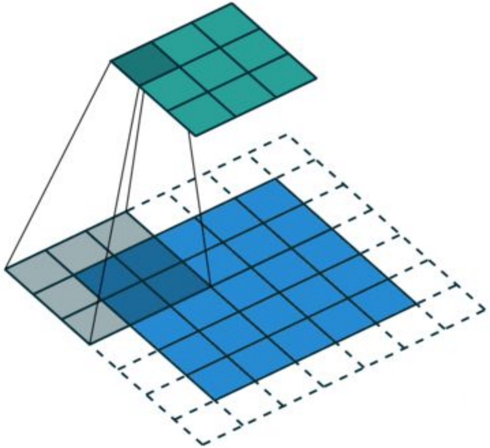

Convolutions

Kernel (filter): small matrix that we use to convolve an image

Convolution:

An operation that "blends" one function with another.

A filter applies a **convolution** operation on an image

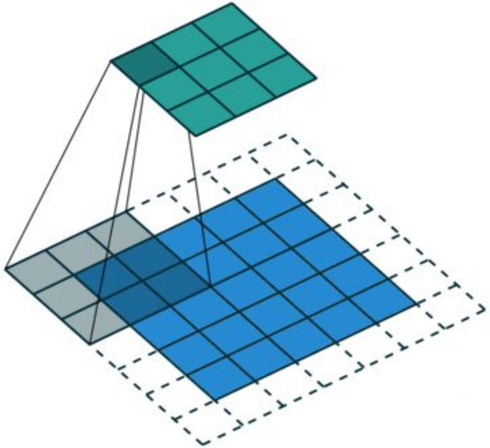

Convolutions Compute

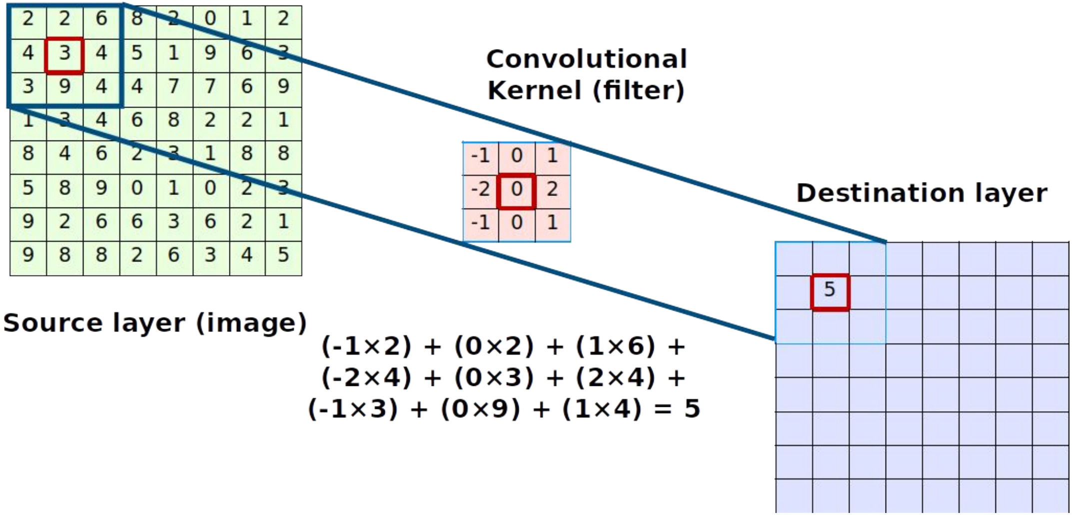

$$(-1\times2) + (0\times2) + (1\times6) + (-2\times4) + (0\times3) + (2\times4) + (-1\times3) + (0\times9) + (1\times4) = 5$$

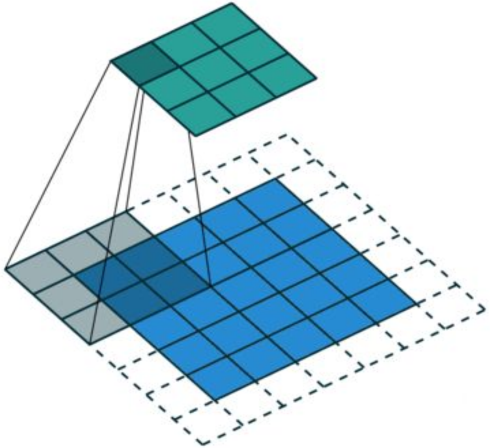

Convolutions Compute

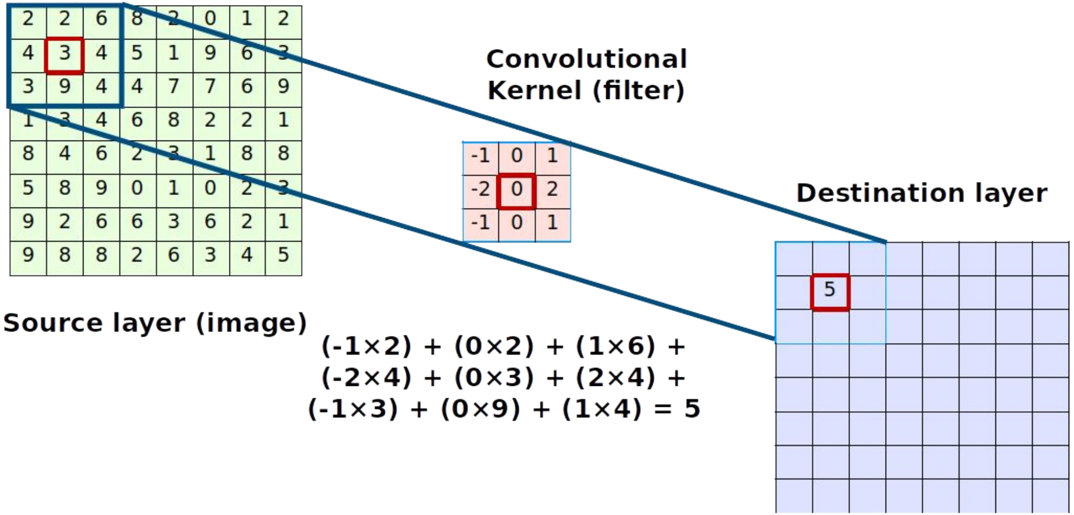

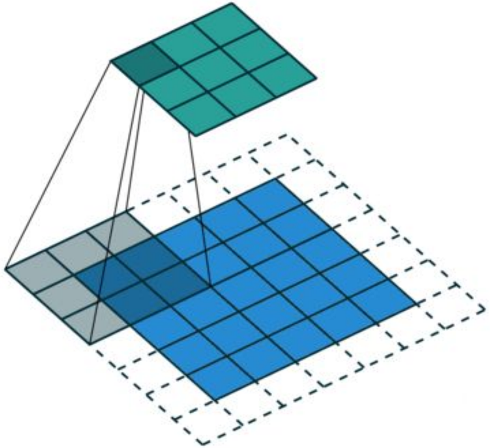

The kernel is **shifted** across the image and produces a point value
The step size in which it shifts is the **stride**
The output is always smaller, we use **padding** to preserve dimensions

Convolutions Compute

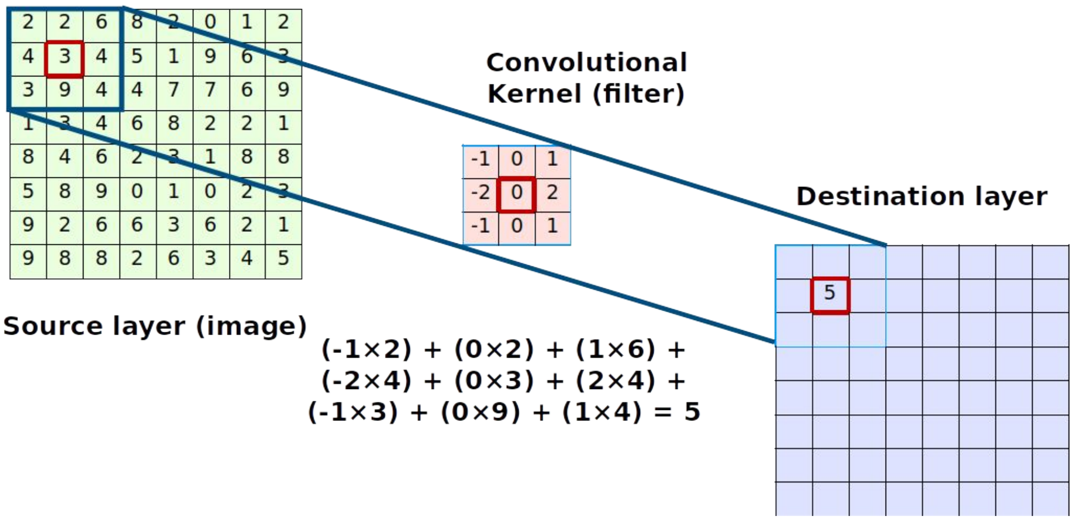

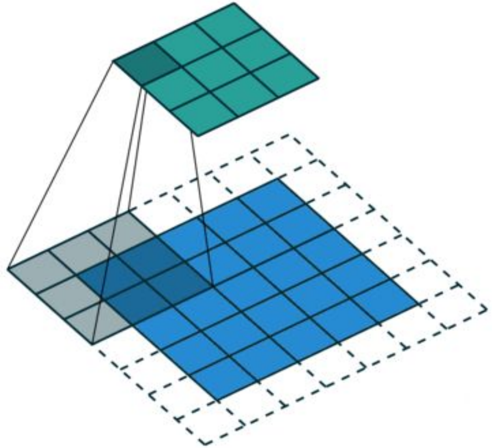

The kernel is **shifted** across the image and produces a point value
The step size in which it shifts is the **stride**
The output is always smaller, we use **padding** to preserve dimensions

The image shrinks according to

Output_size = inputSize -(KernelSize - 1)

Convolutions

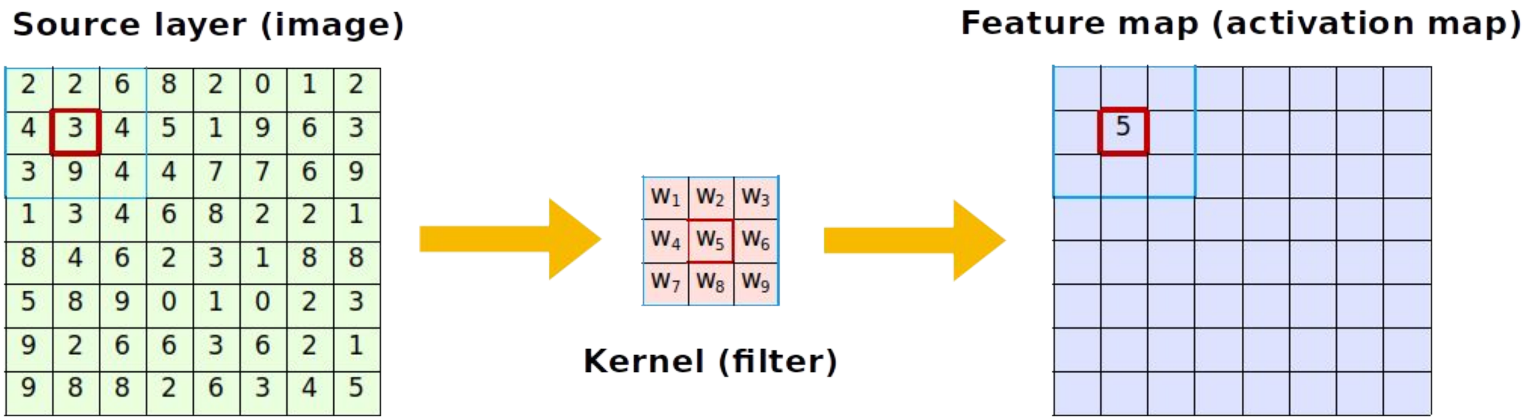

Convolutions

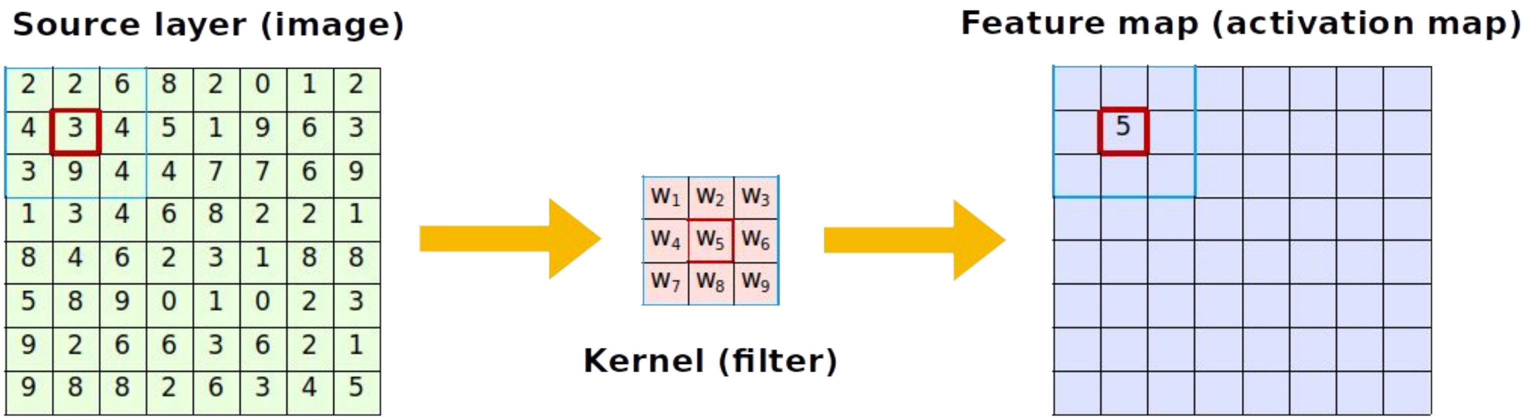

Convolutions

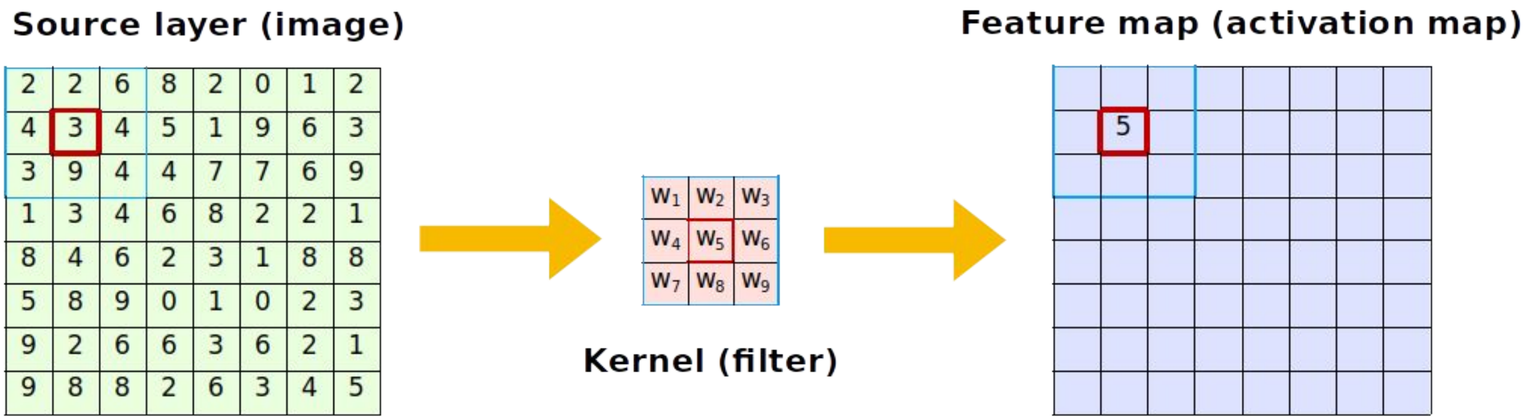

How do we know which kernels to use?

Kernels are learnt and initialized randomly
during training the CNN learns spatial features

Convolutions

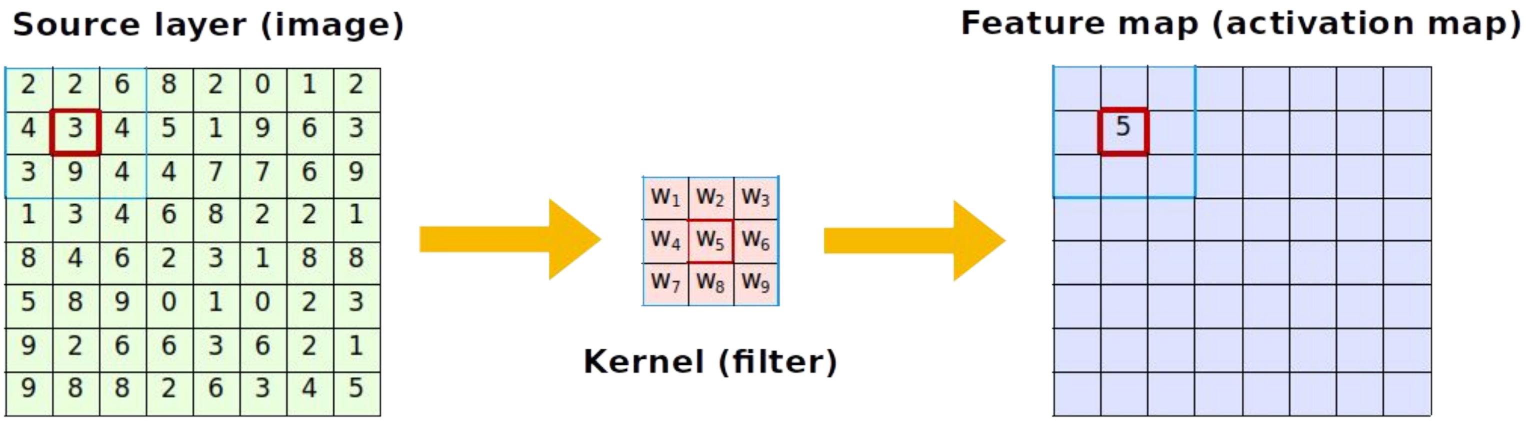

How do we know which kernels to use?

Fully connected

$$\sum_{i\in\text{image}}^{W\times H\times C} \mathbf{x}_i \mathbf{w}_i$$

Kernels are learnt and initialized randomly
during training the CNN learns spatial features

Convolutions

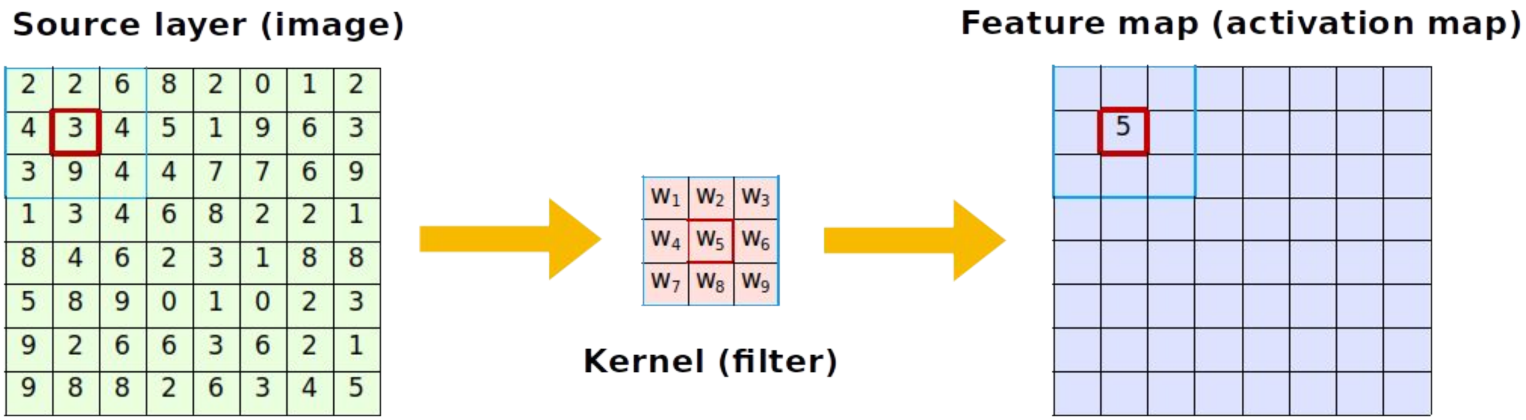

How do we know which kernels to use?

Kernels are learnt and initialized randomly
during training the CNN learns spatial features

Fully connected

$$\sum_{i \in \text{image}}^{W \times H \times C} \mathbf{x}_i \mathbf{w}_i \longrightarrow$$

SURF

Convolutions

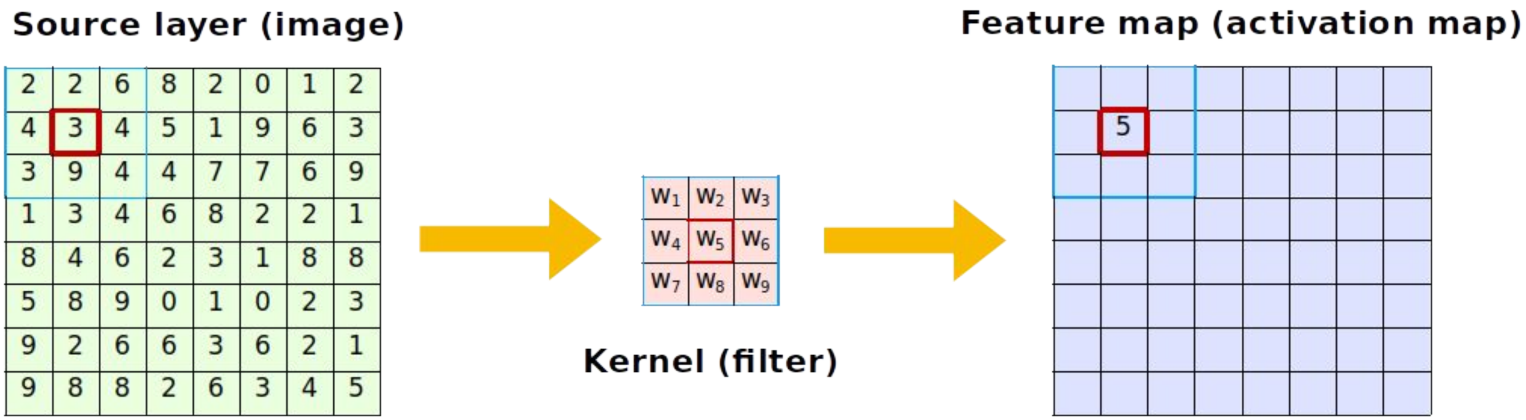

How do we know which kernels to use?

Kernels are learnt and initialized randomly during training the CNN learns spatial features

Fully connected

$$\sum_{i \in \text{image}}^{W \times H \times C} \mathbf{x}_i \mathbf{w}_i \longrightarrow \sum_{i \in 3\times 3}^{W \times H \times C} \mathbf{x}_i \mathbf{w}_i$$

Locally connected shared weights

Pooling and dilated convolutions

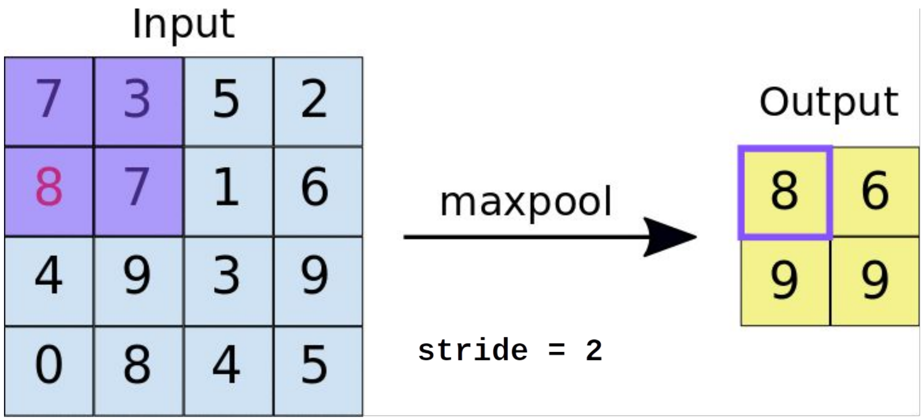

We could discard information gradually by max pooling
Keep the strongest signal, works better than average pooling
Nowadays, we use a bigger stride for dimensionality reduction

Pooling and dilated convolutions

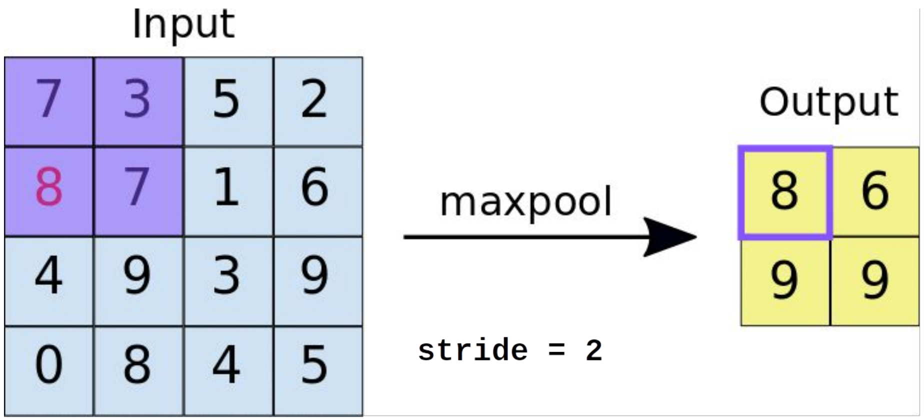

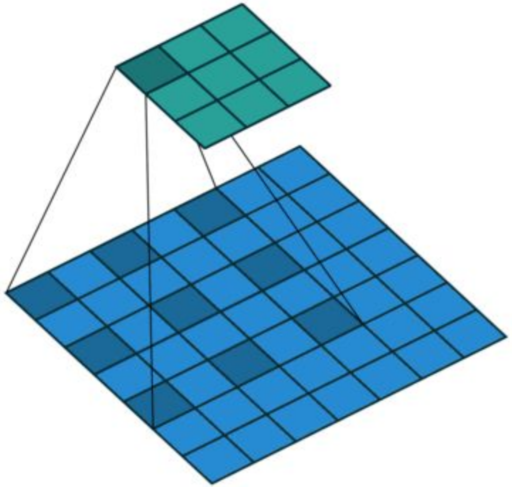

We could discard information gradually by max pooling
Keep the strongest signal, works better than average pooling
Nowadays, we use a bigger stride for dimensionality reduction

Dilated convolutions can be used if you expect your images to have information which are spatially far from each other

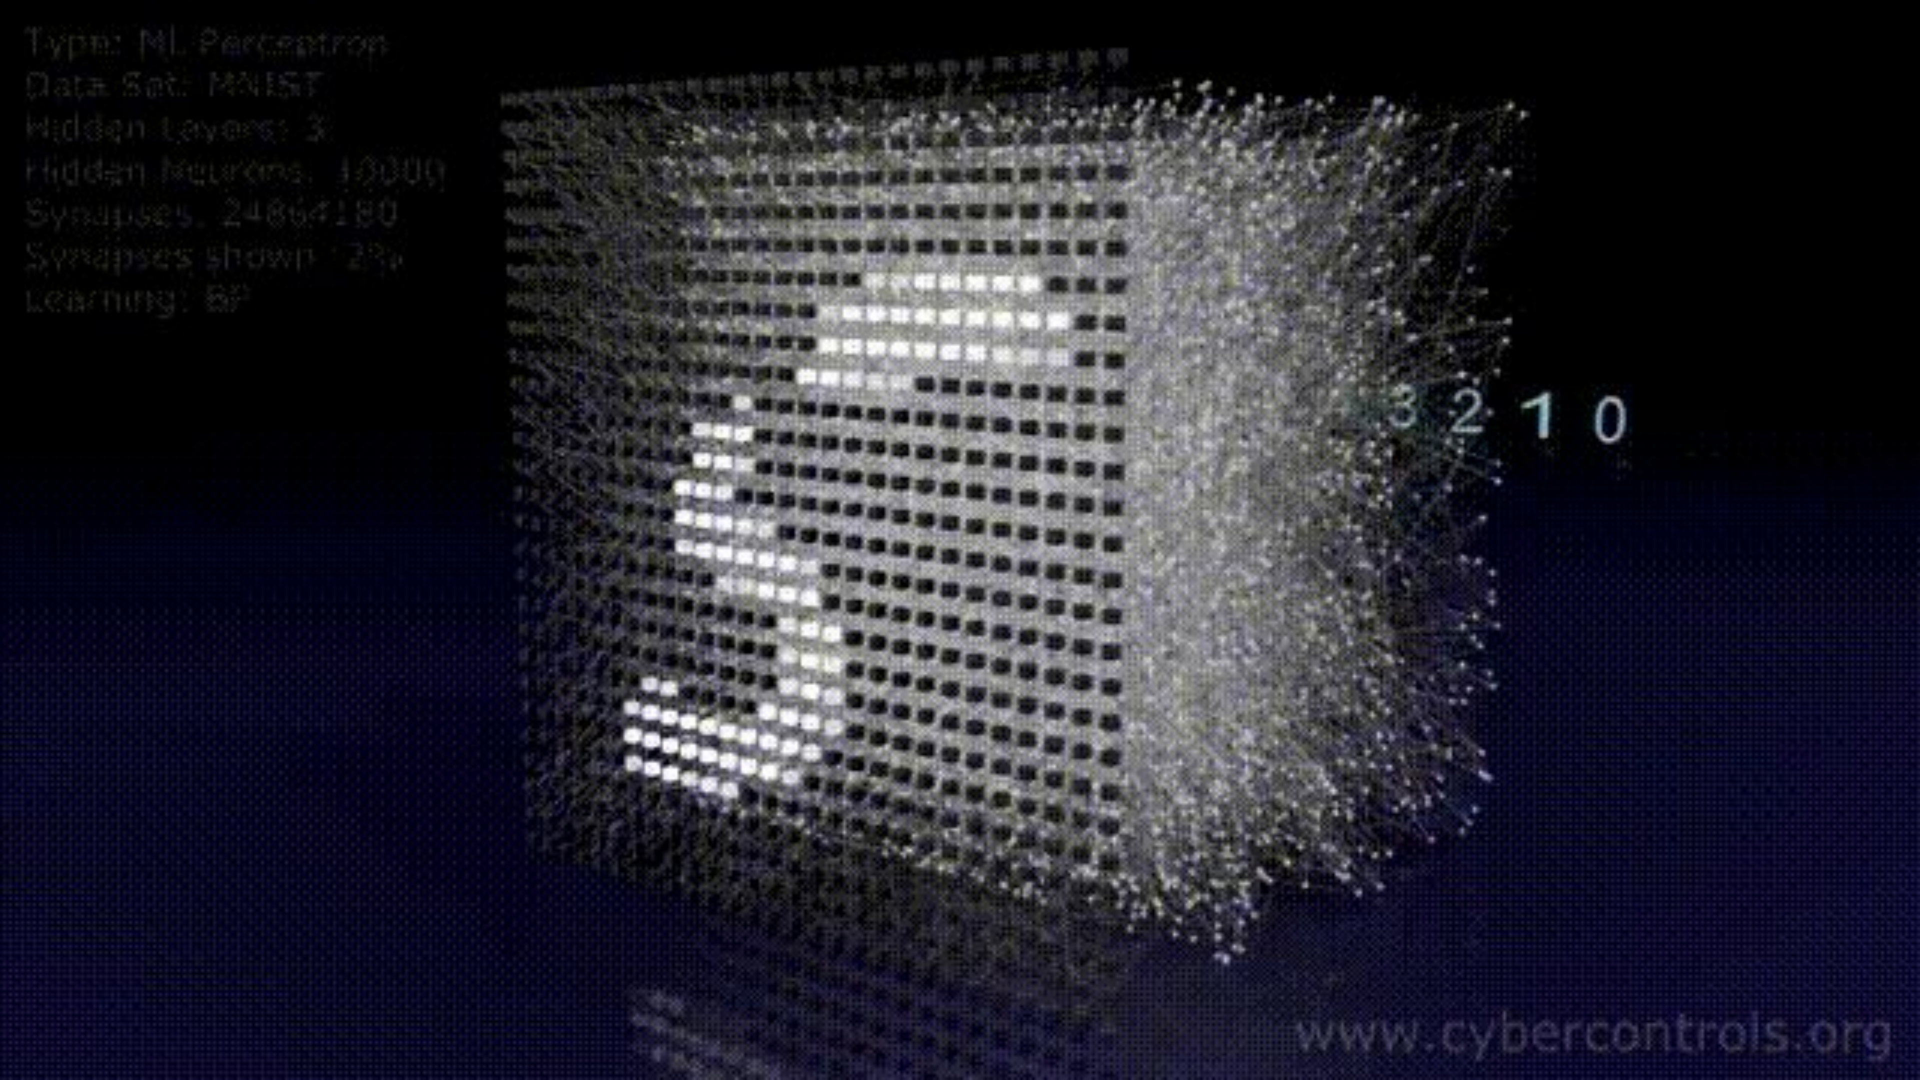
Type: ML Perceptron
Data Set: MNIST
Hidden Layers: 3
Hidden Neurons: 10000
Synapses: 24864180
Learning: BP
3 2 1 0
www.cybercontrols.org

03. CNN Cases

SURF

AlexNet (2012)

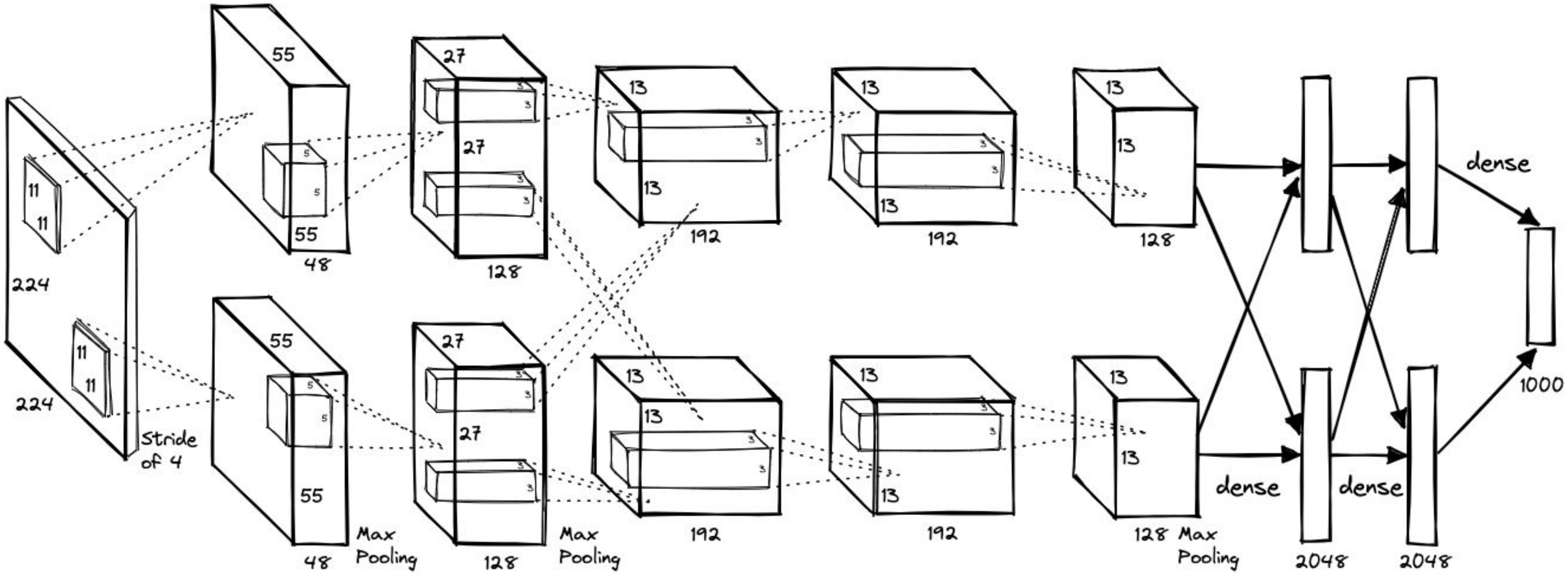

AlexNet (2012)

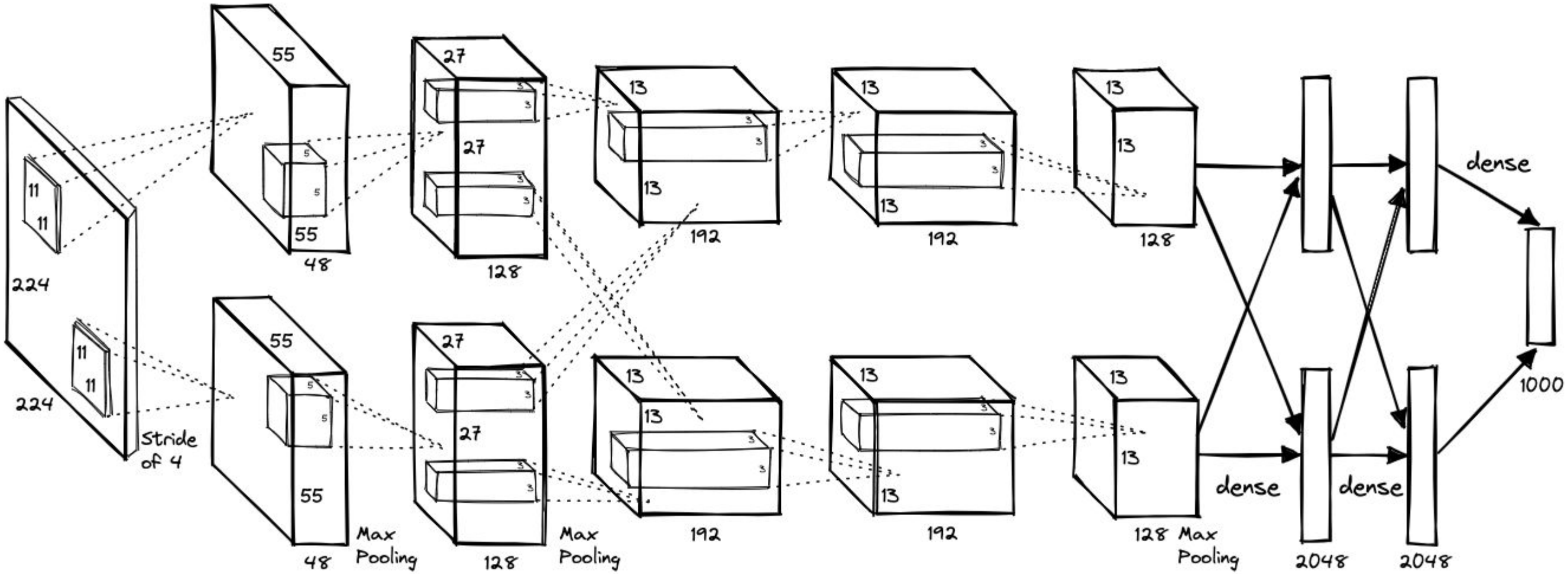

Images become smaller and number of filters increases as we go deeper

8 layers with dropout and ReLU activations trained for 6 days on 2 GPUs

VGG16 and ResNet (2014, 2015)

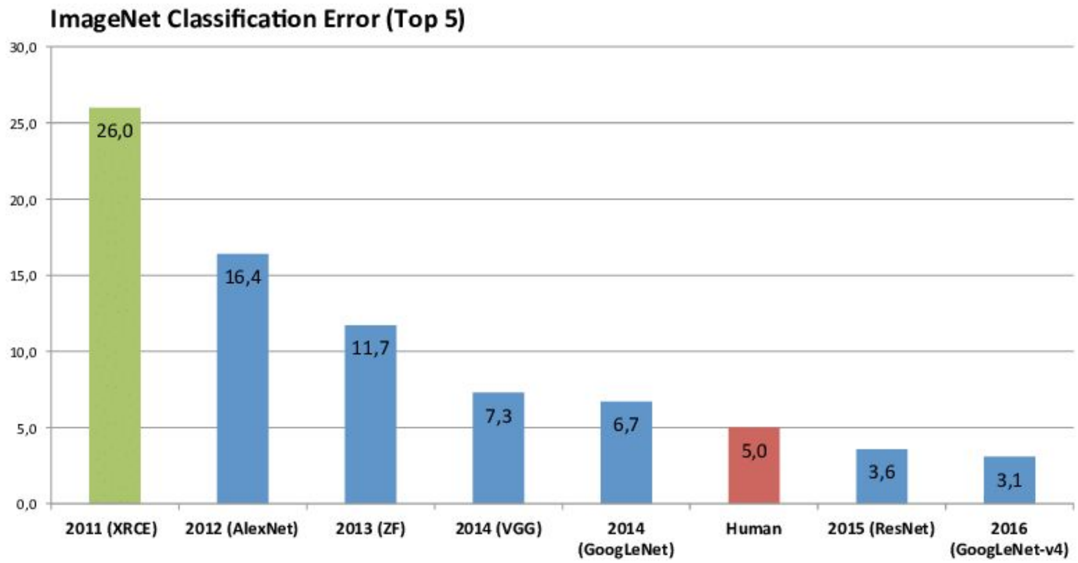

Images become smaller and number of filters increases as we go deeper

8 layers with dropout and ReLU activations trained for 6 days on 2 GPUs

VGG16 and ResNet (2014, 2015)

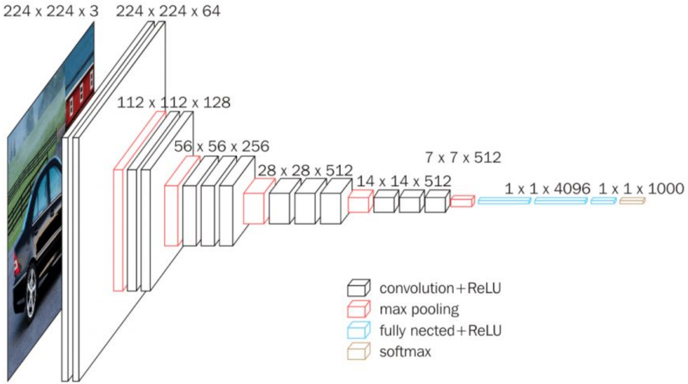

VGG16: construct large and deep models (120M params)

Vanishing gradient problem

VGG16 and ResNet (2014, 2015)

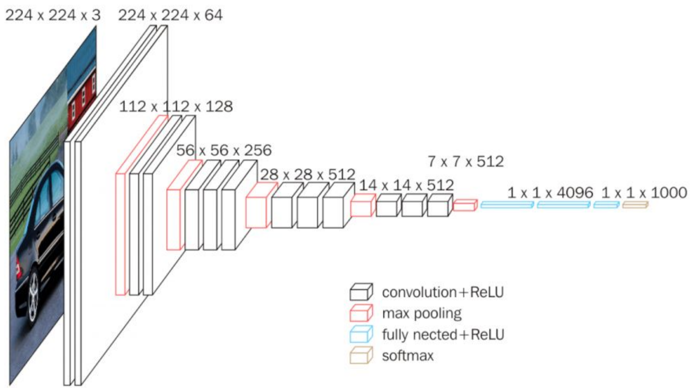

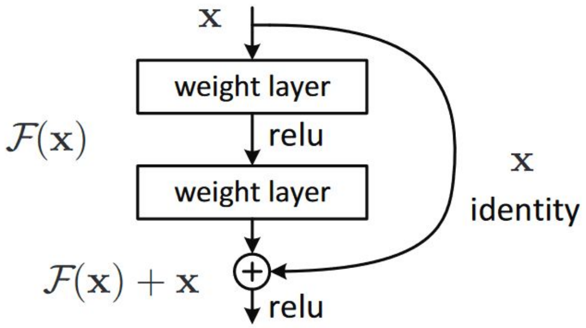

VGG16: construct large and deep models (120M params)

Vanishing gradient problem

ResNet (2015)

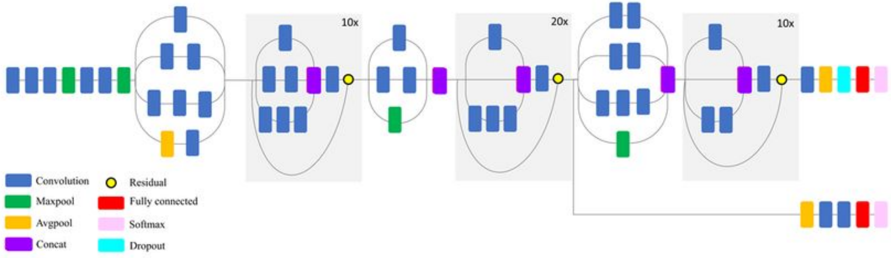
10x
20x
10x
Convolution
Maxpool
Avgpool
Concat
Residual
Fully connected
Softmax
Dropout

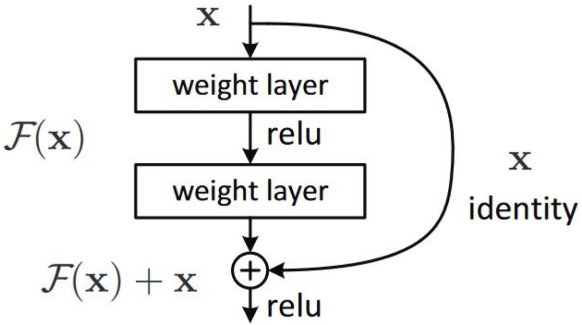
$\mathbf{x}$
weight layer
relu
weight layer
$\mathcal{F}(\mathbf{x})$
$\mathbf{x}$
identity
$\mathcal{F}(\mathbf{x}) + \mathbf{x}$
relu

ResNet (2015)

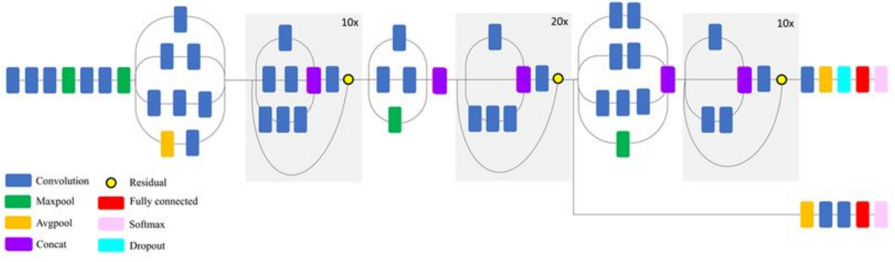

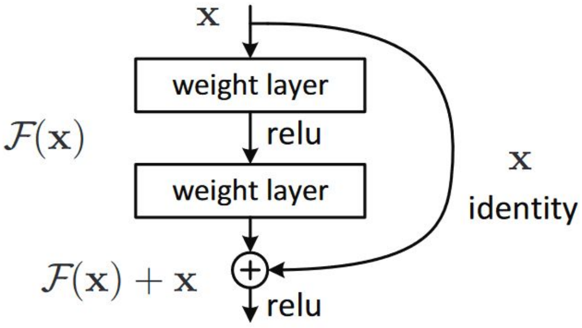

ResNets: construct large and deep models with skip connections, able to train up to 152 layers (!)

Higher abstraction and less nuisance from vanishing or exploding gradients

ResNet (2015)

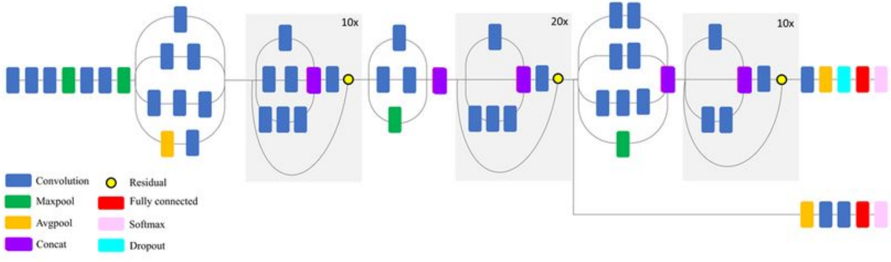

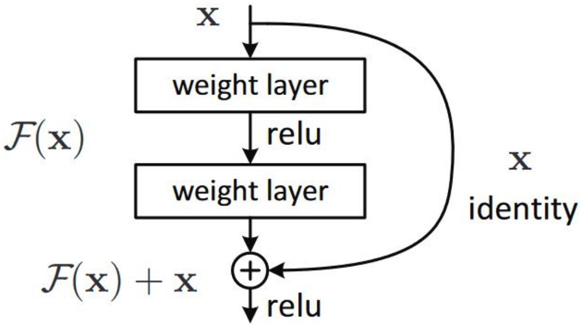

ResNets: construct large and deep models with skip connections, able to train up to 152 layers (!)

Higher abstraction and less nuisance from vanishing or exploding gradients

Early layers learn features that get progressively more abstract, we can **preserve** and **control** the flow of information with **skip connections**

DenseNet (2016)

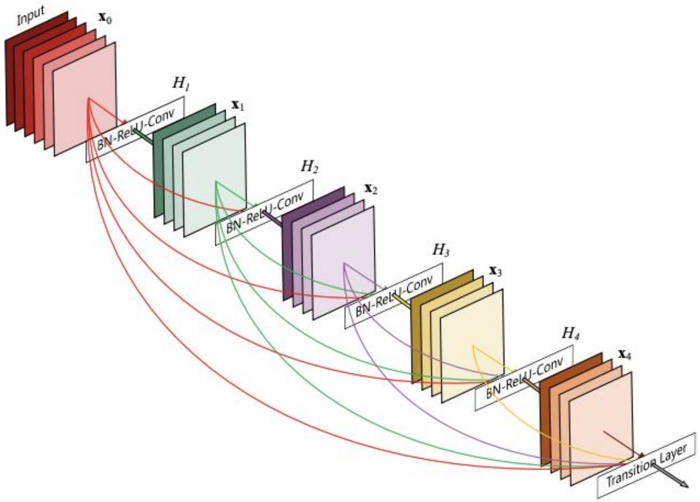

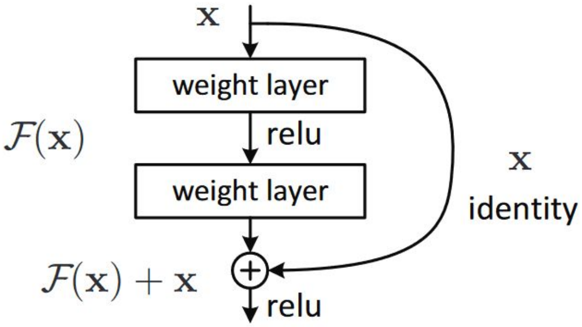

ResNets: construct large and deep models with skip connections, able to train up to 152 layers (!)

Higher abstraction and less nuisance from vanishing or exploding gradients

Early layers learn features that get progressively more abstract, we can **preserve** and **control** the flow of information with **skip connections**

SURF

DenseNet (2016)

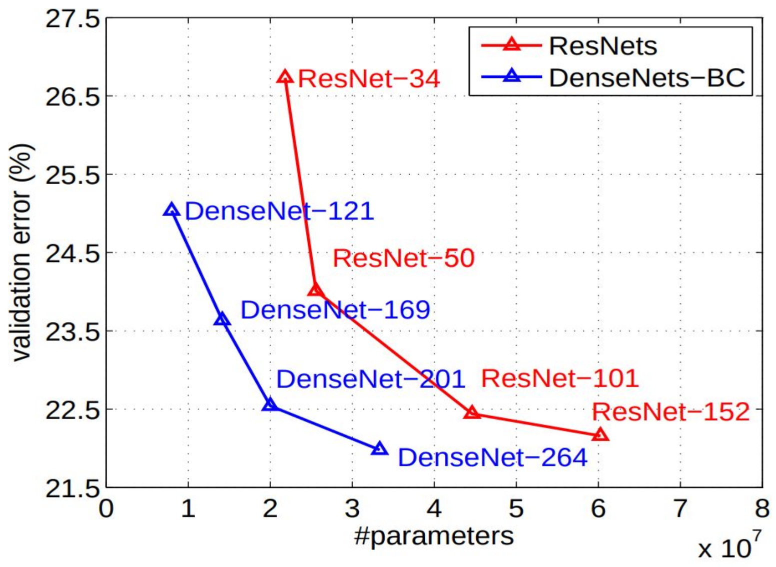

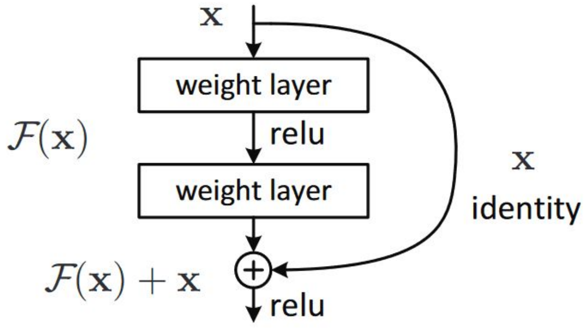

ResNets: construct large and deep models with skip connections, able to train up to 152 layers (!)

Higher abstraction and less nuisance from vanishing or exploding gradients

Early layers learn features that get progressively more abstract, we can **preserve** and **control** the flow of information with **skip connections**

Transfer Learning

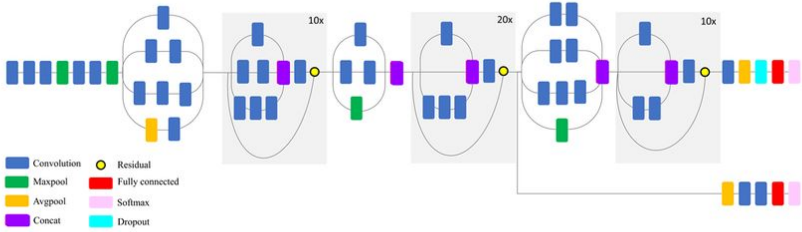

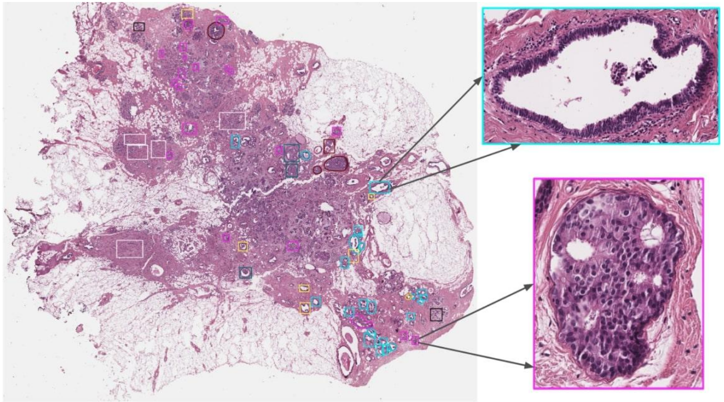

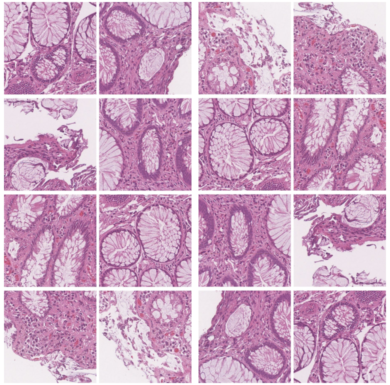

Thank You

High Performance Machine Learning Group